

# Chapter 4: Threads & Concurrency



# Chapter 4: Threads

---

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples



# Objectives

---

- Identify the *basic components of a thread*, and *contrast threads and processes*
- Describe the *benefits* and *challenges* of designing *multithreaded applications*
- Illustrate *different approaches to implicit threading* including thread pools, fork-join, and Grand Central Dispatch
- Describe how the Windows and Linux operating systems represent threads
- Design multithreaded applications using the Pthreads, Java, and Windows threading APIs


# Motivation

---


- Most modern applications are *multithreaded*
- *Threads run within application*
- Multiple tasks within the application can be implemented by separate threads
  - Update display
  - Fetch data
  - Spell checking
  - Answer a network request
- Process creation is heavy-weight while *thread creation is light-weight*
- Can simplify code, increase efficiency
- Kernels are generally multithreaded



# Single and Multithreaded Processes




single-threaded process



multithreaded process



# Multithreaded Server Architecture



# Benefits

---

- *Responsiveness* – may allow continued execution if part of process is blocked, especially important for user interfaces
- *Resource Sharing* – threads share resources of process, easier than shared memory or message passing (IPC)
- *Economy* – cheaper than process creation, thread switching lower overhead than context switching
- *Scalability* – process can take advantage of multicore architectures

# Multicore Programming


---

- *Multicore* or *multiprocessor systems* putting pressure on programmers, challenges include:
  - Dividing activities
  - Balance
  - Data splitting
  - Data dependency
  - Testing and debugging
- *Parallelism* implies a system can perform more than one task simultaneously
- *Concurrency* supports more than one task making progress
  - Single processor / core, scheduler providing concurrency




# Concurrency vs. Parallelism

## □ Concurrent execution on single-core system:



## □ Parallelism on a multi-core system:




# Multicore Programming

---

## □ Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- **Task parallelism** – distributing threads across cores, each thread performing unique operation

# Data and Task Parallelism




# Amdahl's Law

---

- Identifies *performance gains* from adding additional cores to an application that has both serial and parallel components
  - $s$  is serial portion
  - $n$  processing cores
  - That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
  - As  $n$  approaches infinity, speedup approaches  $1/s$
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores
- But does the law take into account *contemporary multicore systems*?



# Amdahl's Law



# User Threads and Kernel Threads


---

- *User threads* - management done by user-level threads library
- Three primary thread libraries:
  - POSIX Pthreads
  - Windows threads
  - Java threads
- *Kernel threads* - supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
  - Windows, Linux, Mac OS X
  - iOS, Android



# User and Kernel Threads

---



# Multithreading Models


---

- ❑ Many-to-One
- ❑ One-to-One
- ❑ Many-to-Many

# Many-to-One

---


- ❑ *Many user-level threads* mapped to *single kernel thread*
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- ❑ *Few systems currently use this model*
- ❑ Examples:
  - Solaris Green Threads
  - GNU Portable Threads



# One-to-One


---

- ❑ *Each user-level thread* maps to *one kernel thread*
- ❑ Creating a user-level thread creates a kernel thread
- ❑ More concurrency than many-to-one
- ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Examples
  - Windows
  - Linux



# Many-to-Many Model

- ❑ Allows *many user level threads* to be mapped to *many kernel threads*
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Windows with the **ThreadFiber** package
- ❑ Otherwise *not very common*



# Thread Libraries

---

- ❑ **Thread library** provides programmer with API for creating and managing threads
- ❑ Two primary ways of implementing
  - Library entirely *in user space*
  - Kernel-level library *supported by the OS*

# Pthreads

---

- May be provided either as *user-level* or *kernel-level*
- A **POSIX standard (IEEE 1003.1c) API** for thread creation and synchronization
- Specification, not implementation
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)

# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}
```



# Pthreads Example (cont)

---

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



# Implicit Threading

---

- ❑ Growing in popularity as numbers of threads increase, program correctness more difficult with *explicit threads*
- ❑ *Creation and management of threads done by compilers and run-time libraries* rather than programmers
- ❑ Five methods explored
  - Thread Pools
  - Fork-Join
  - OpenMP
  - Grand Central Dispatch
  - Intel Threading Building Blocks



# Thread Pools


---

- Create a *number of threads in a pool* where they await work
- Advantages:
  - Usually slightly *faster* to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be *bound* to the size of the pool
  - Separating task to be performed from mechanics of creating task allows different strategies for running task
    - ▶ i.e., Tasks could be scheduled to run periodically



# Fork-Join Parallelism

- *Multiple threads (tasks) are forked, and then joined.*



# Fork-Join Parallelism

---

- General algorithm for *fork-join strategy*:

```
Task(problem)
    if problem is small enough
        solve the problem directly
    else
        subtask1 = fork(new Task(subset of problem)
        subtask2 = fork(new Task(subset of problem)

        result1 = join(subtask1)
        result2 = join(subtask2)

    return combined results
```



# Threading Issues

---

- ❑ Semantics of `fork()` and `exec()` system calls
- ❑ Signal handling
  - Synchronous and asynchronous
- ❑ Thread cancellation of target thread
  - Asynchronous or deferred
- ❑ Thread-local storage
- ❑ Scheduler Activations



# Semantics of fork() and exec()

---

- ❑ Does **fork()** duplicate only the calling thread or all threads?
  - Some UNIXes have two versions of fork
- ❑ **exec()** usually works as normal – replace the running process including all threads

# Signal Handling

---

- ❑ **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- ❑ A *signal handler* is used to process signals
  - Signal is generated by particular event
  - Signal is delivered to a process
  - Signal is handled by one of two signal handlers
    - ▶ default
    - ▶ user-defined
- ❑ Every signal has a default handler that kernel runs when handling signal
  - User-defined signal handler can override default
  - For single-threaded, signal delivered to process



# Signal Handling (Cont.)

---

- *Where* should a signal be delivered for *multi-threaded*?
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the process

# Thread Cancellation

---

- Terminating a thread before it has finished
- Thread to be canceled is target thread

- Two general approaches:
  - Asynchronous cancellation terminates the target thread immediately
  - Deferred cancellation allows the target thread to periodically check if it should be cancelled
- **Pthread** code to create and cancel a thread:

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);  
  
/* wait for the thread to terminate */  
pthread_join(tid,NULL);
```



# Thread Cancellation (Cont.)

---

- ❑ Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state


Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- ❑ If thread has cancellation disabled, cancellation remains pending until thread enables it
- ❑ Default type is deferred
  - Cancellation only occurs when thread reaches cancellation point
    - i.e., `pthread_testcancel()`
    - Then cleanup handler is invoked
- ❑ On Linux systems, thread cancellation is handled through signals



# Scheduler Activations

- ❑ Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ❑ Typically use an intermediate data structure between user and kernel threads – *lightweight process (LWP)*
  - Appears to be a *virtual processor* on which process can schedule user thread to run
  - Each LWP attached to kernel thread
  - How many LWPs to create?
- ❑ Scheduler activations provide upcalls - a communication mechanism from the kernel to the upcall handler in the thread library
- ❑ This communication allows an application to maintain the correct number kernel threads



# End of Chapter 4



*What Is an*  
**OPERATING SYSTEM (OS)**  
*and How Does It Work*

CLEVERISM.COM