

Chapter 10.A:

File-System Interface



Chapter 10.A: Outline

- File Concept
- Access Methods
- Disk and Directory Structure
- File-System Mounting
- File Sharing
- Protection



Objectives

- To explain the *function of file systems*
- To describe the *interfaces to file systems*
- To discuss *file-system design tradeoffs*, including access methods, file sharing, file locking, and directory structures
- To explore *file-system protection*

■ Contiguous *logical address space*

■ Types:

- Data

- ▶ complex
- ▶ numeric
- ▶ character
- ▶ binary

- Program

■ Contents defined by file's creator

- Many types

- ▶ Text file
- ▶ Source file
- ▶ Executable file



- *Name* – only information kept in human-readable form
- *Identifier* – unique tag (number) identifies file within file system
- *Type* – needed for systems that support different types
- *Location* – pointer to file location on device
- *Size* – current file size
- *Protection* – controls who can do reading, writing, executing
- *Time, date, and user identification* – data for protection, security, and usage monitoring
- Information about files are kept in the directory structure, which is maintained on the disk
- Many variations, including extended file attributes such as file checksum



A window of file info on Mac OS X



- File is an *abstract data type*
 - *Create*
 - *Write* – at write pointer location
 - *Read* – at read pointer location
 - *Reposition* within file (or seek)
 - *Delete*
 - *Truncate*
- *Open(F_i)* – search the directory structure on disk for entry F_i , and move the content of entry to memory
- *Close(F_i)* – move the content of entry F_i in memory to directory structure on disk



- Several pieces of data are needed to manage open files:
 - *Open-file table*: tracks open files
 - *File pointer*: pointer to last read/write location, per process that has the file open
 - *File-open count*: counter of number of times a file is open – to allow removal of data from open-file table when last processes closes it
 - *Disk location of the file*: cache of data access information
 - *Access rights*: per-process access mode information



- Provided by some operating systems and file systems
 - Similar to *reader-writer locks*
 - *Shared lock* similar to reader lock – several processes can acquire concurrently
 - *Exclusive lock* similar to writer lock
- Mediates access to a file
- Mandatory or advisory:
 - *Mandatory* – access is denied depending on locks held and requested
 - *Advisory* – processes can find status of locks and decide what to do



```
import java.io.*;
import java.nio.channels.*;
public class LockingExample {
    public static final boolean EXCLUSIVE = false;
    public static final boolean SHARED = true;
    public static void main(String arsg[]) throws IOException {
        FileLock sharedLock = null;
        FileLock exclusiveLock = null;
        try {
            RandomAccessFile raf = new RandomAccessFile("file.txt",
"rw");
            // get the channel for the file
            FileChannel ch = raf.getChannel();
            // this locks the first half of the file - exclusive
            exclusiveLock = ch.lock(0, raf.length()/2, EXCLUSIVE);
            /** Now modify the data . . . */
            // release the lock
            exclusiveLock.release();
        }
    }
}
```





File Locking Example – Java API (Cont.)

```
        // this locks the second half of the file - shared
        sharedLock = ch.lock(raf.length()/2+1,
raf.length(), SHARED);

        /** Now read the data . . . */

        // release the lock
        sharedLock.release();

    } catch (java.io.IOException ioe) {
        System.err.println(ioe);
    }finally {
        if (exclusiveLock != null)
            exclusiveLock.release();
        if (sharedLock != null)
            sharedLock.release();
    }
}
}
```

File Types – Name, Extension

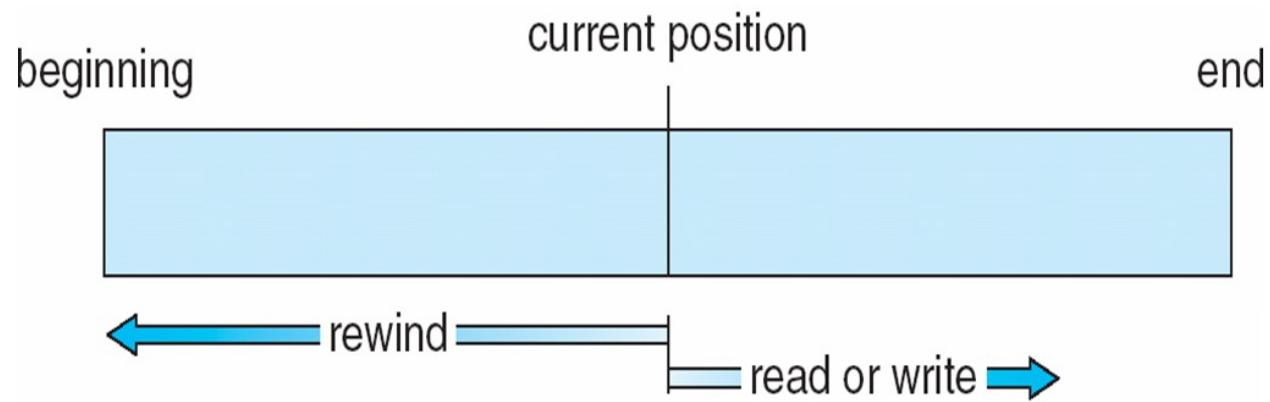
file type	usual extension	function
executable	exe, com, bin or none	ready-to-run machine-language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rtf, doc	various word-processor formats
library	lib, a, so, dll	libraries of routines for programmers
print or view	ps, pdf, jpg	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes compressed, for archiving or storage
multimedia	mpeg, mov, rm, mp3, avi	binary file containing audio or A/V information



- *None* - sequence of words or bytes
- *Simple record structures*
 - Lines
 - Fixed length
 - Variable length
- *Complex structures*
 - Formatted document
 - Relocatable load file (i.e., executable file)
- Can simulate last two with first method by inserting appropriate control characters
- Who decides:
 - Operating system
 - Program



Sequential-Access File



■ *Sequential Access*

read next

write next

reset

no read after last write
(rewrite)

■ *Direct Access* – file is *fixed-length logical records*

read n

write n

position to n

read next

write next

rewrite n

n = *relative block number*

- Relative block numbers allow OS to decide where file should be placed
- See *allocation problem* in Chapter 12





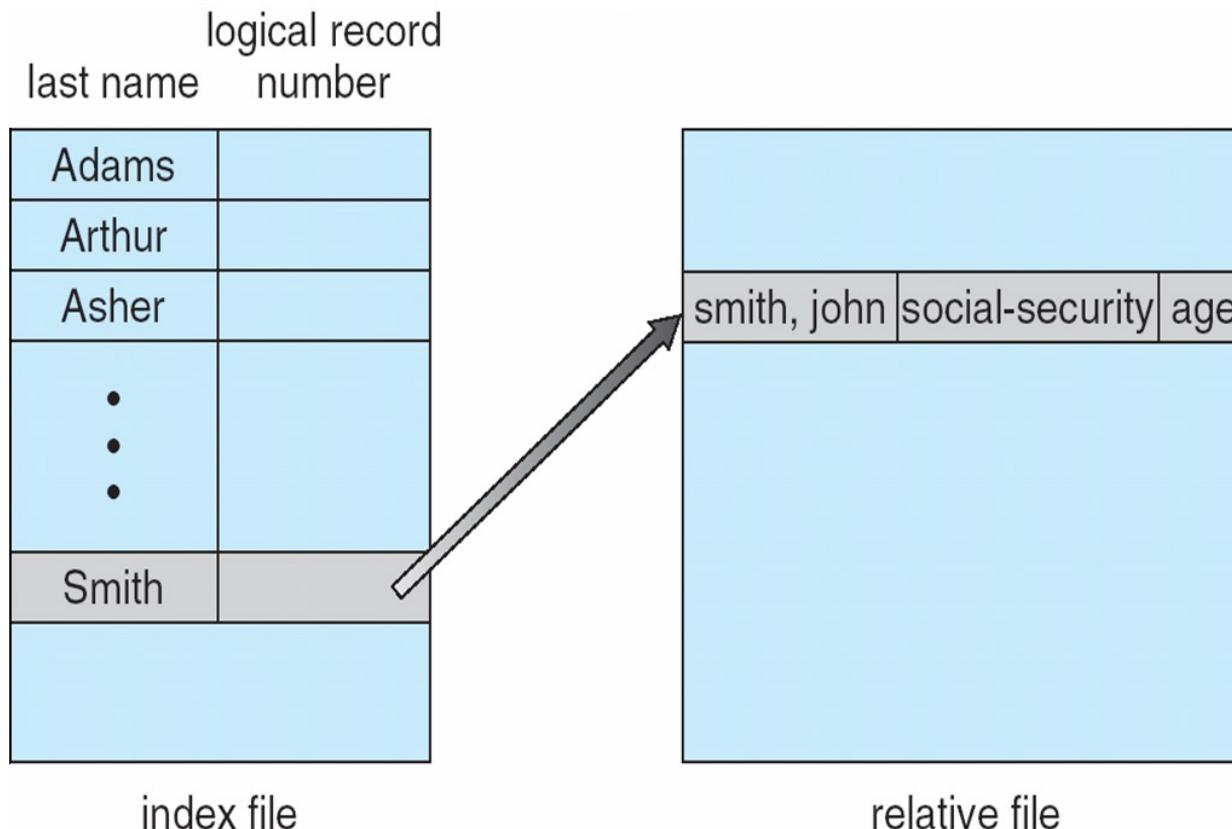
Simulation of Sequential-Access on Direct-Access File

sequential access	implementation for direct access
<i>reset</i>	$cp = 0;$
<i>read next</i>	$read cp;$ $cp = cp + 1;$
<i>write next</i>	$write cp;$ $cp = cp + 1;$

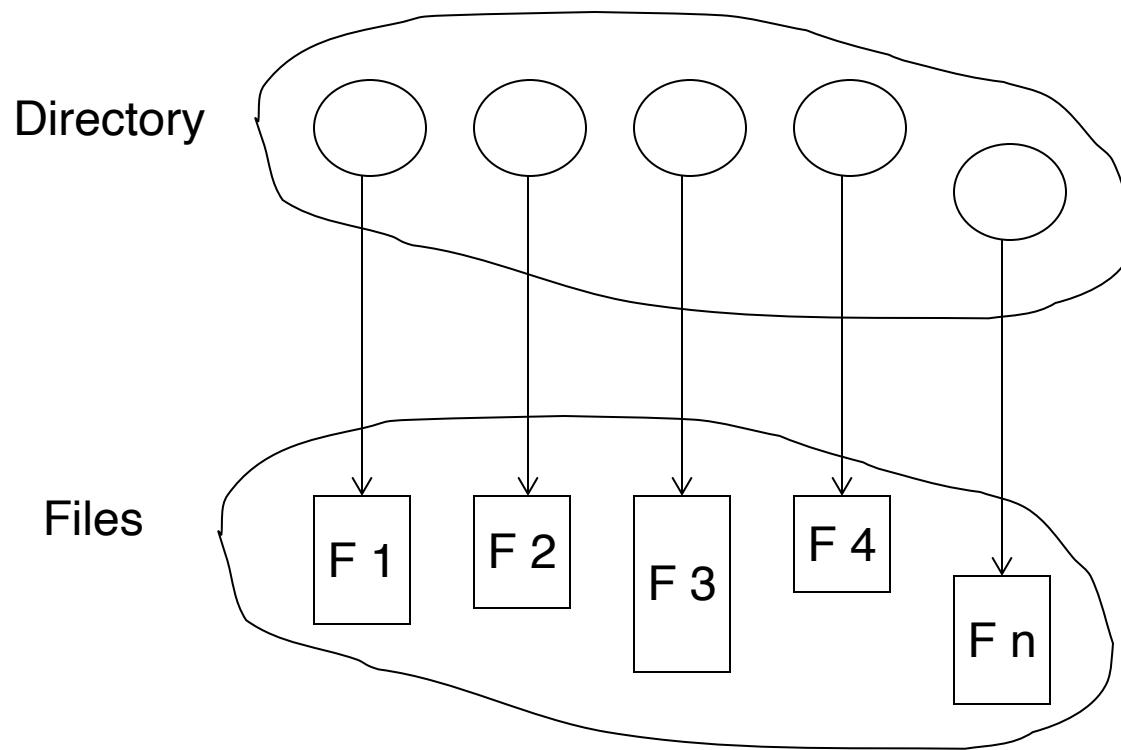
- Can be built on top of base methods
- General involve creation of an *index* for the file
- Keep index in memory for fast determination of location of data to be operated on (consider UPC code plus record of data about that item)
- If too large, index (in memory) of the index (on disk)
- E.g., **IBM *Indexed Sequential-Access Method (ISAM)***
 - Small master index, points to disk blocks of secondary index
 - File kept sorted on a defined key
 - All done by the OS
- **VMS** operating system provides index and relative files as another example (see next slide)



Example of Index and Relative Files



- A *collection of nodes* containing information about all files



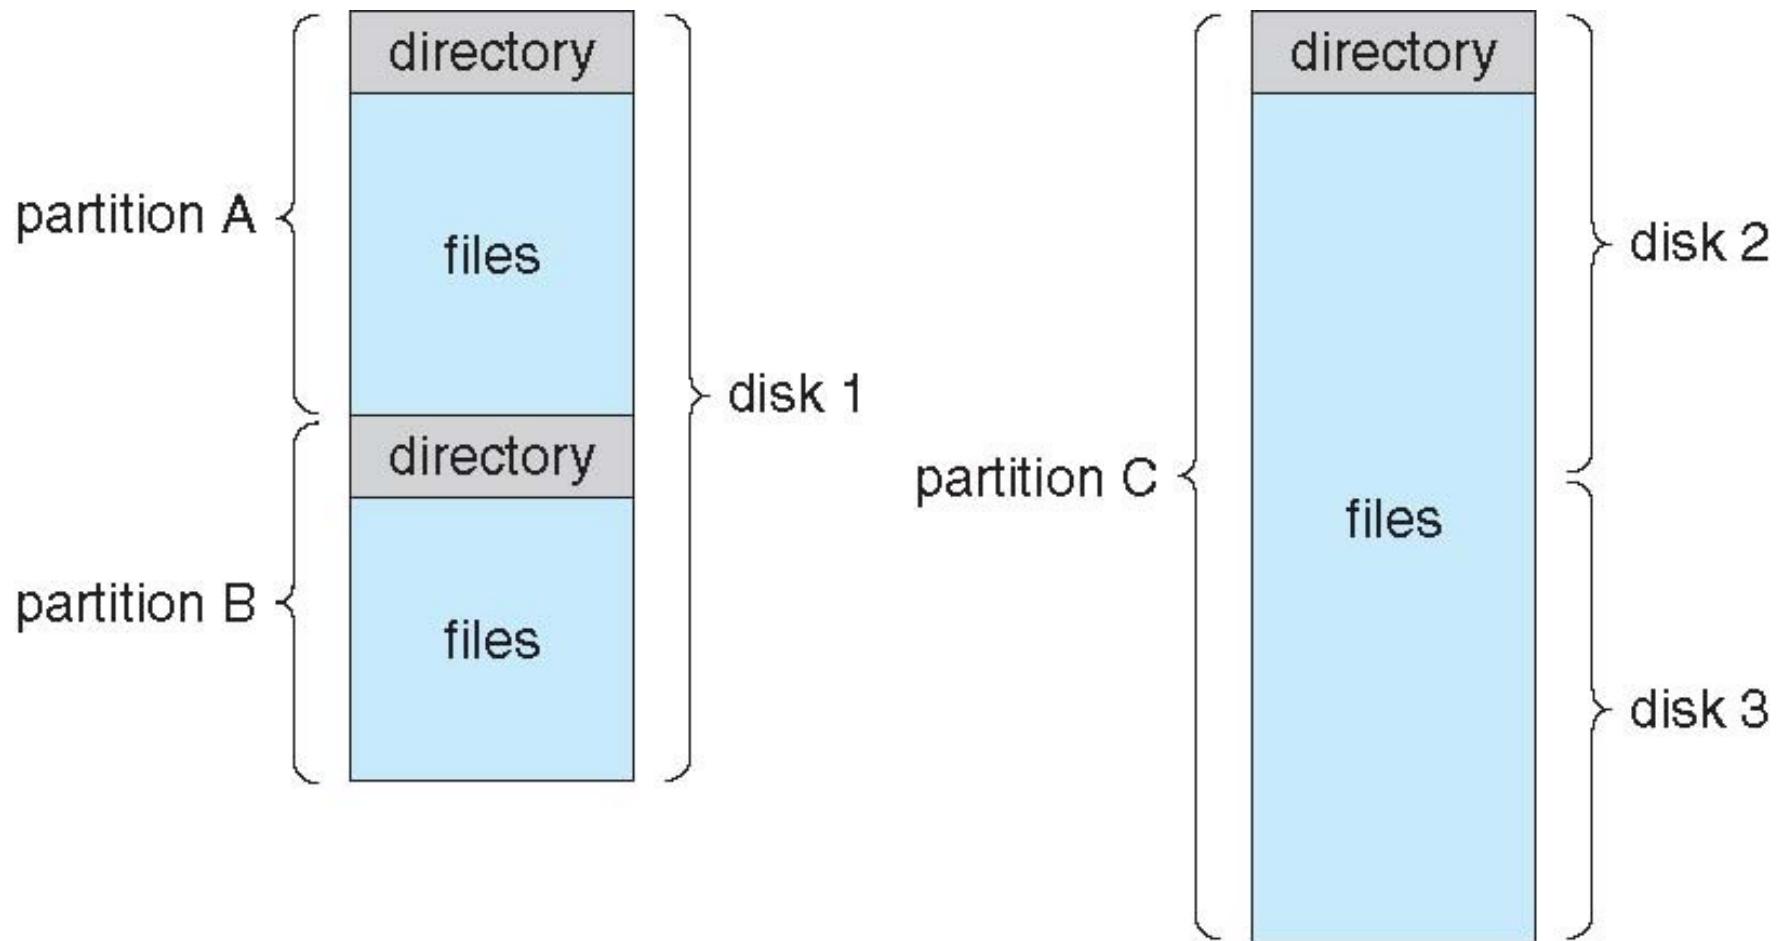
Both the directory structure and the files reside on disk



- *Disk* can be subdivided into *partitions*
 - Disks or partitions can be **RAID** protected against failure
 - Disk or partition can be used *raw* – without a file system, or *formatted* with a file system
- Partitions also known as minidisks, slices
- Entity containing file system known as a *volume*
 - Each volume containing file system also tracks that file system's info in device directory or volume table of contents
- As well as general-purpose file systems there are many special-purpose file systems, frequently all within the same operating system or computer



A Typical File-System Organization



Types of File Systems

- We mostly talk of *general-purpose file systems*
- But systems frequently have many file systems, some general- and some special- purpose
- E.g., **Solaris** has
 - *tmpfs* – memory-based volatile FS for fast, temporary I/O
 - *objfs* – interface into kernel memory to get kernel symbols for debugging
 - *ctfs* – contract file system for managing daemons
 - *lofs* – loopback file system allows one FS to be accessed in place of another
 - *procfs* – kernel interface to process structures
 - *ufs, zfs* – general purpose file systems

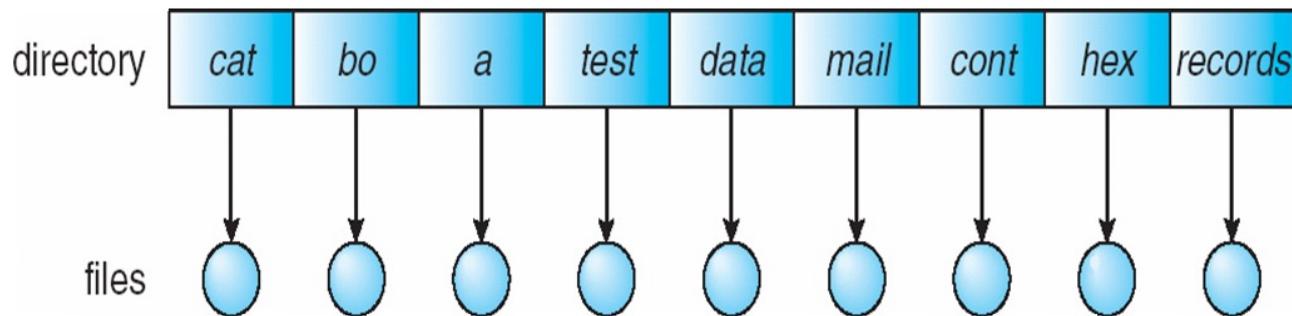


- *Search* for a file
- *Create* a file
- *Delete* a file
- *List* a directory
- *Rename* a file
- *Traverse* the file system

- The directory is organized logically to obtain
 - *Efficiency* – locating a file quickly
 - *Naming* – convenient to users
 - ▶ Two users can have same name for different files
 - ▶ The same file can have several different names
 - *Grouping* – logical grouping of files by properties, (e.g., all Java programs, all games, ...)

Single-Level Directory

- *A single directory for all users*

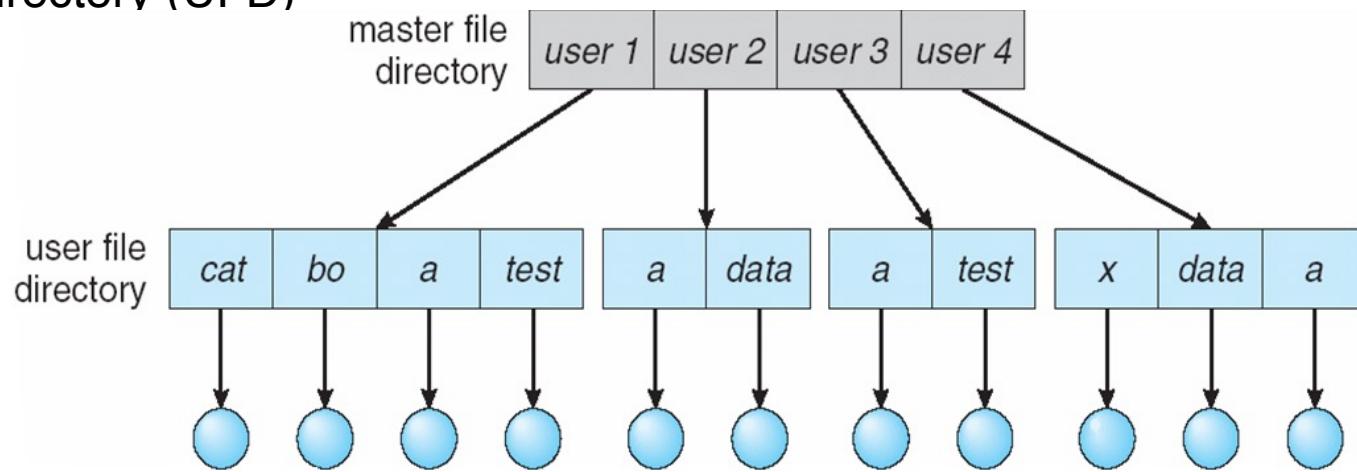


- Naming problem
- Grouping problem

Two-Level Directory

- *Separate directory for each user*

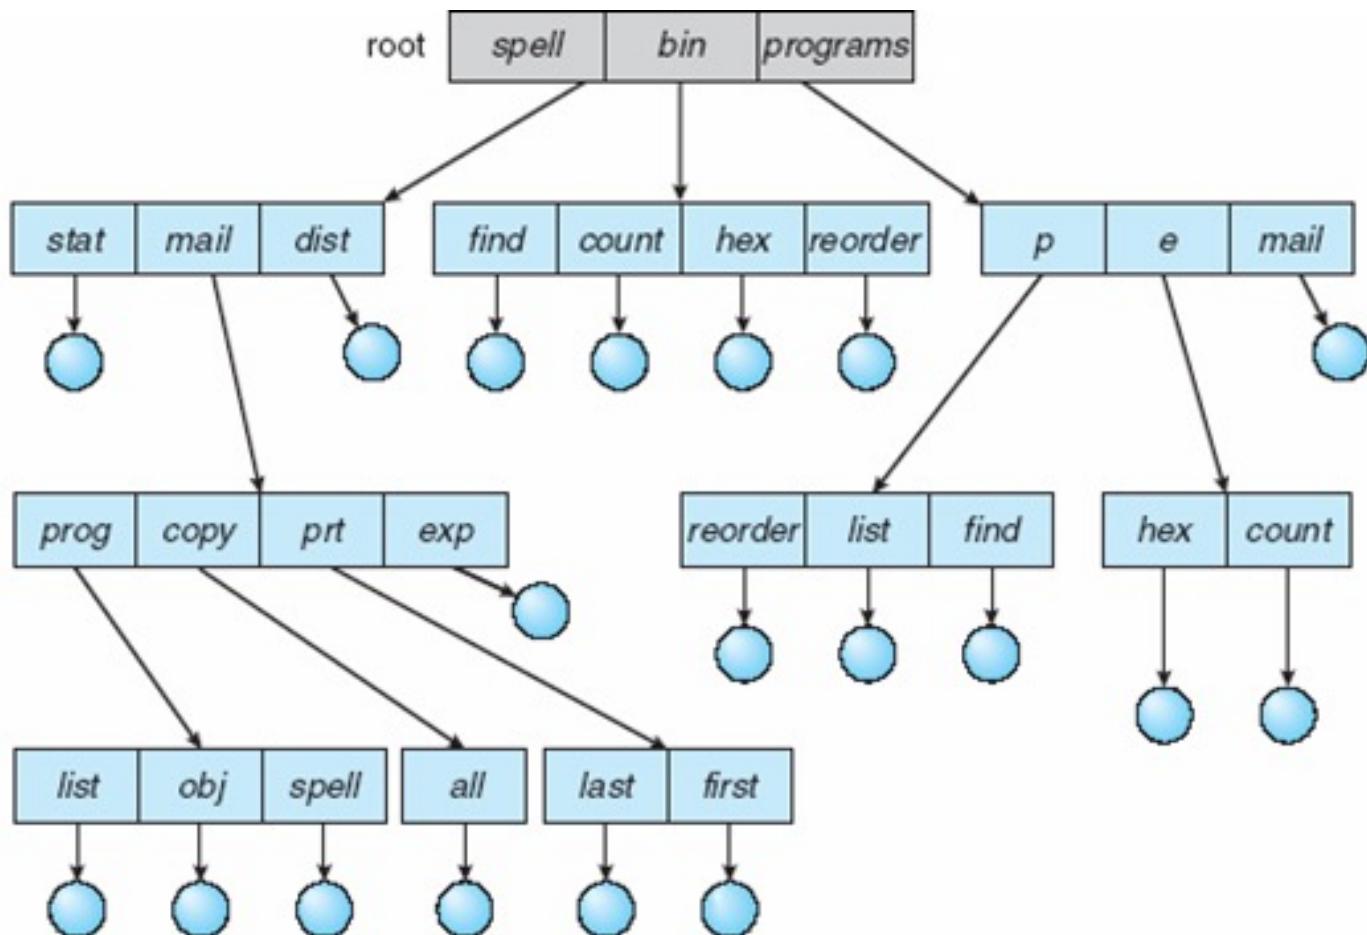
- Master file directory (MFD)
- User file directory (UFD)



- Path name
- Can have the same file name for different user
- Efficient searching
- No grouping capability



Tree-Structured Directories





Tree-Structured Directories (Cont.)

- Efficient searching
- Grouping Capability
- Current directory (or working directory)

- E.g., For Linux OS,

```
cd /spell/mail/proc
```

```
type list
```



- Using *absolute* or *relative* path name
- Creating a new file is done in current directory
- Delete a file

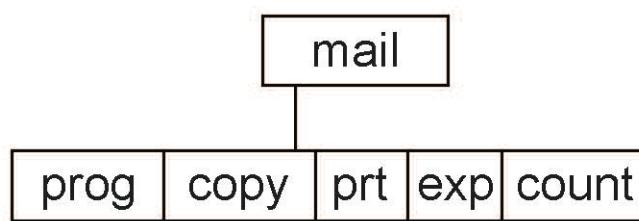
```
rm <file-name>
```

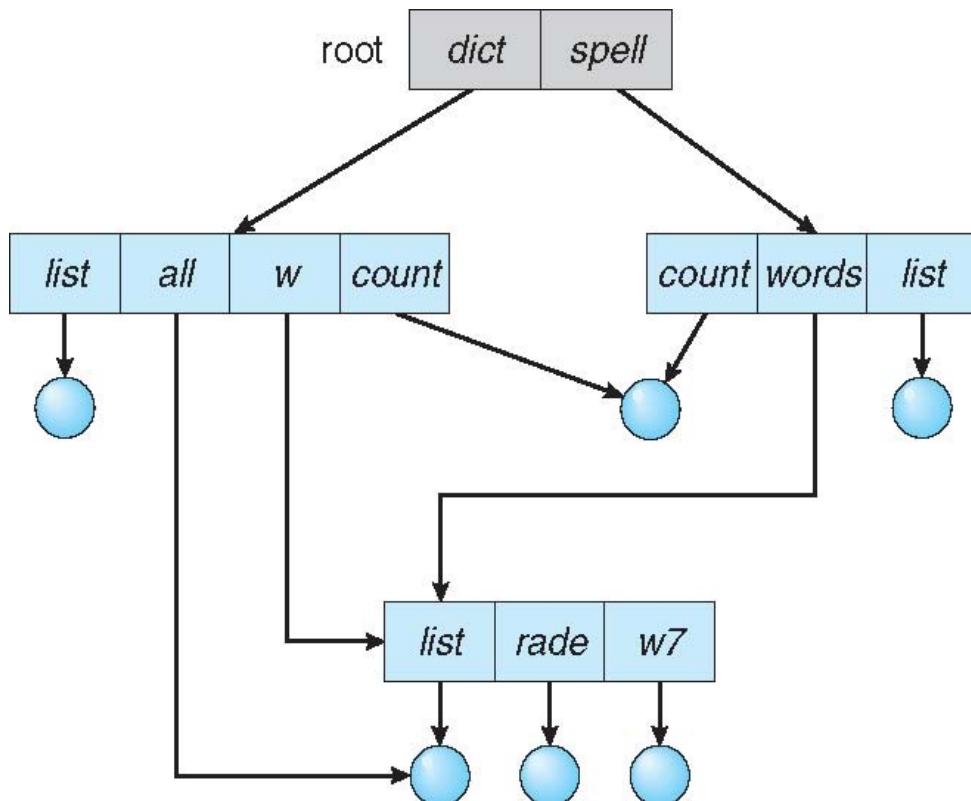
- Creating a new subdirectory is done in current directory

```
mkdir <dir-name>
```

- Example: if in current directory /mail

```
mkdir count
```



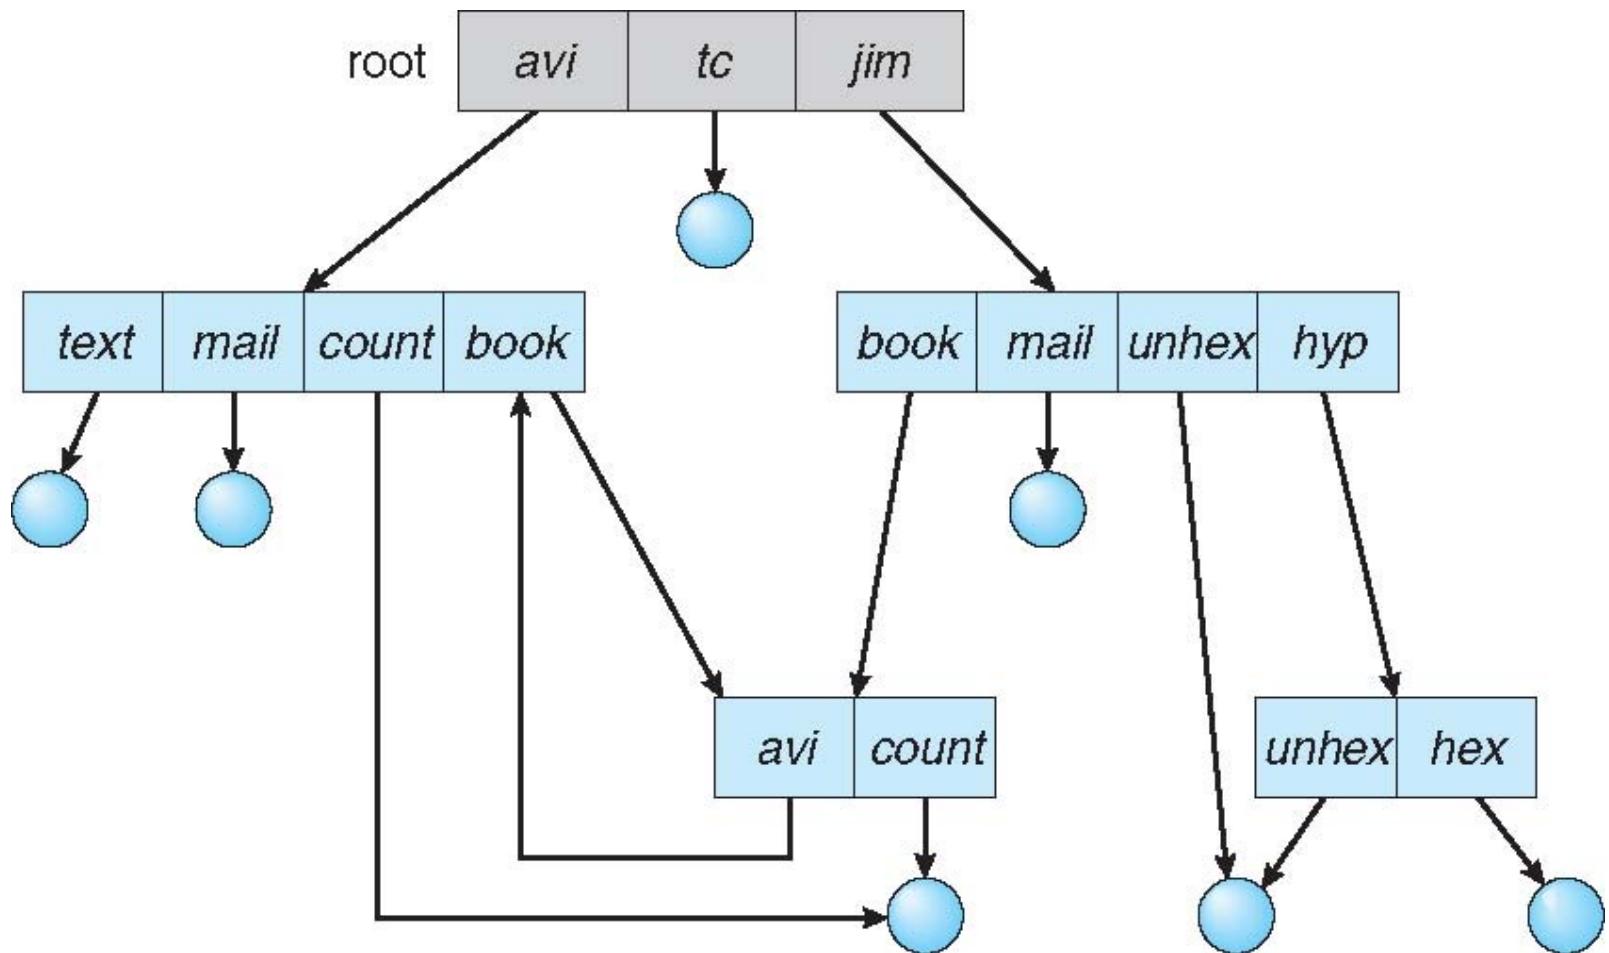


- Have shared subdirectories and files
 - Two different names (*aliasing*)
 - ▶ If *dict* deletes *list* \Rightarrow dangling pointer.
 - Solutions:
 - ▶ *Backpointers*, so we can delete all pointers
 - ▶ Variable size records a problem
 - Backpointers using a daisy chain organization
 - ▶ *Entry-hold-count* solution

■ New directory entry type

- *Link* – another name (pointer) to an existing file
- *Resolve the link* – follow pointer to locate the file

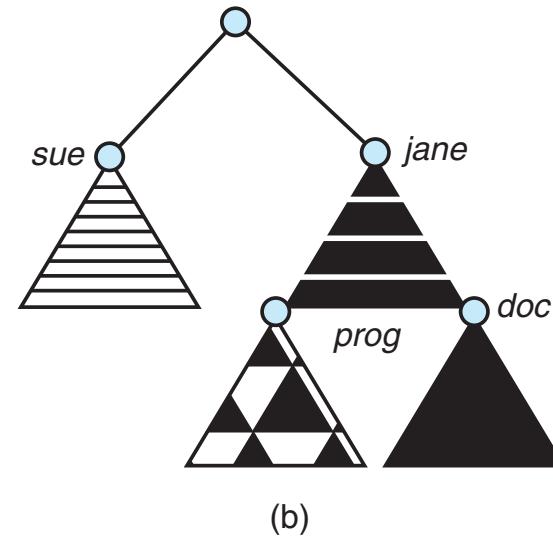
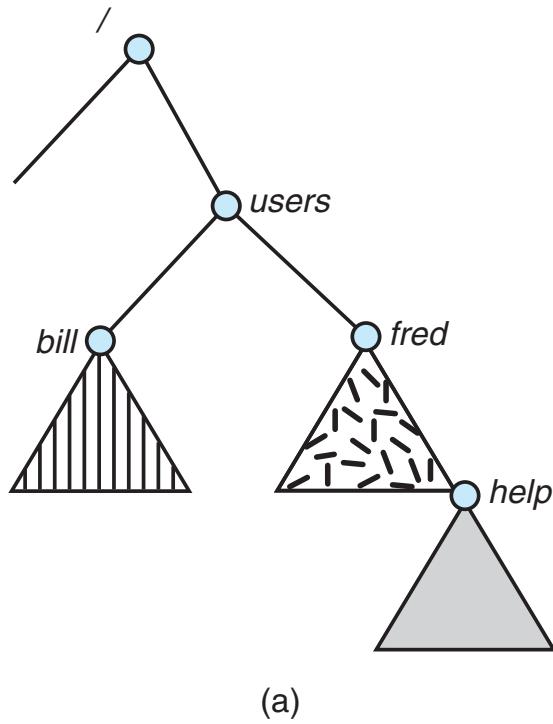




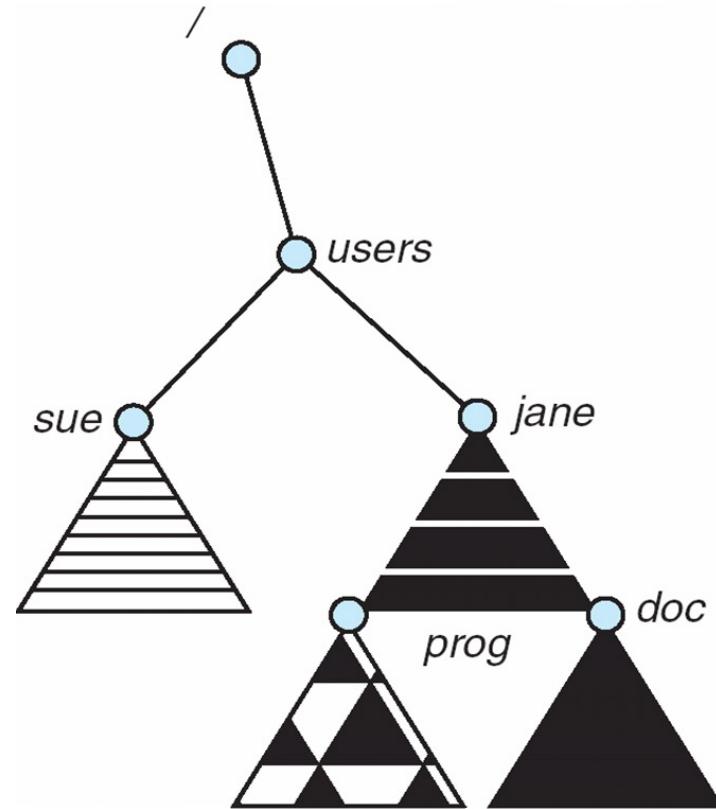
■ How do we guarantee no cycles?

- Allow *only links to file* not subdirectories
- *Garbage collection*
- Every time a new link is added use a *cycle detection algorithm* to determine whether it is OK

- A file system must be *mounted* before it can be *accessed*
- A unmounted file system (i.e., Fig. (b)) is mounted at a mount point



Mount Point



- Sharing of files on *multi-user systems* is desirable
- Sharing may be done through a protection scheme
- On distributed systems, files may be shared across a network
 - *Network File System (NFS)* is a common distributed file-sharing method
- If multi-user system
 - Owner of a file / directory
 - ▶ *User IDs* identify users, allowing permissions and protections to be per-user
 - Group of a file / directory
 - ▶ *Group IDs* allow users to be in groups, permitting group access rights



- Uses *networking* to allow file system access between systems
 - Manually via programs like FTP
 - Automatically, seamlessly using *distributed file systems*
 - Semi automatically via the world wide web
- *Client-server model* allows clients to mount remote file systems from servers
 - Server can serve multiple clients
 - Client and user-on-client identification is insecure or complicated
 - **NFS** is standard UNIX client-server file sharing protocol
 - **CIFS** is standard Windows protocol
 - Standard operating system file calls are translated into remote calls
- *Distributed Information Systems* (distributed naming services) such as **LDAP**, **DNS**, **NIS**, **Active Directory** implement unified access to information needed for remote computing



- All file systems have *failure modes*
 - For example corruption of directory structures or other non-user data, called *metadata*
- Remote file systems add new failure modes, due to *network failure*, *server failure*
- Recovery from failure can involve state information about status of each remote request
- Stateless protocols such as **NFS v.3** include all information in each request, allowing easy recovery but less security



- Specify how multiple users are to access a shared file simultaneously
 - Similar to Ch. 5 *process synchronization algorithms*
 - ▶ Tend to be less complex due to disk I/O and network latency (for remote file systems)
 - *Andrew File System (AFS)* implemented complex remote file sharing semantics
 - ▶ AFS has session semantics
 - Writes only visible to sessions starting after the file is closed
 - *Unix file system (UFS)* implements:
 - ▶ Writes to an open file visible immediately to other users of the same open file
 - ▶ Sharing file pointer to allow multiple users to read and write concurrently



■ *File owner/creator* should be able to control:

- what can be done
- by whom

■ *Types of access*

- Read
- Write
- Execute
- Append
- Delete
- List



- Mode of access: *read* (R), *write* (W), *execute* (X)
- Three classes of users on Unix / Linux

			RWX
owner	group	public	
chmod	761	game	
a) owner access		7	⇒ 1 1 1
			RWX
b) group access		6	⇒ 1 1 0
			RWX
c) public access		1	⇒ 0 0 1

- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.

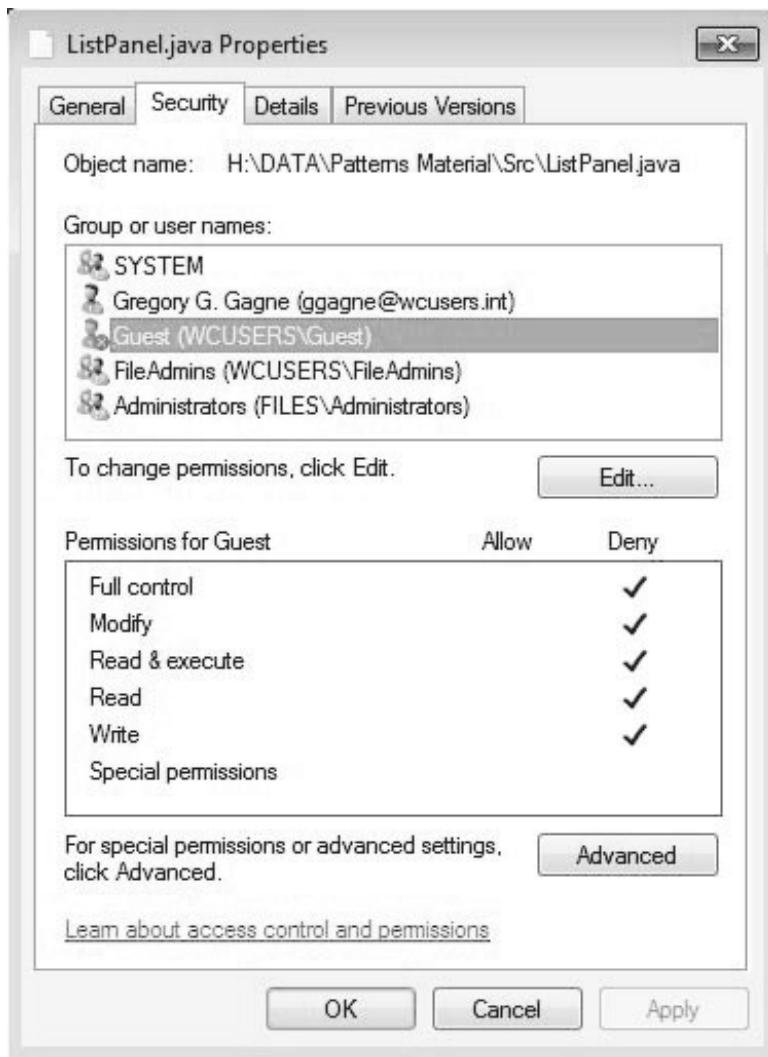
Attach a group to a file

chgrp G game





Windows 7 Access-Control List Management





A Sample UNIX Directory Listing

-rw-rw-r--	1	pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	5	pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	2	pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	2	pbg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	1	pbg	staff	9423	Feb 24 2003	program.c
-rwxr-xr-x	1	pbg	staff	20471	Feb 24 2003	program
drwx--x--x	4	pbg	faculty	512	Jul 31 10:31	lib/
drwx-----	3	pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	3	pbg	staff	512	Jul 8 09:35	test/

- A *file* is an abstract data type defined and implemented by the operating system. It is a *sequence of logical records*. A logical record may be a byte, a line (of fixed or variable length), or a more complex data item. The operating system may specifically support various record types or may leave that support to the application program.
- A major task for the operating system is to *map the logical file concept onto physical storage devices* such as hard disk or NVM device. Since the physical record size of the device may not be the same as the logical record size, it may be necessary to order logical records into physical records. Again, this task may be supported by the operating system or left for the application program.



- Within a *file system*, it is useful to create *directories* to allow files to be organized. A *single-level directory* in a multiuser system causes naming problems, since each file must have a unique name. A *two-level directory* solves this problem by creating a separate directory for each user's files. The directory lists the files by name and includes the file's location on the disk, length, type, owner, time of creation, time of last use, ...
- The natural generalization of a two-level directory is a *tree-structured directory*. A tree-structured directory allows a user to create subdirectories to organize files. *Acyclic-graph directory structures* enable users to share subdirectories and files but complicate searching and deletion. A general graph structure allows complete flexibility in the sharing of files and directories but sometimes requires garbage collection to recover unused disk space.



- *Remote file systems* present challenges in reliability, performance, and security. Distributed information systems maintain user, host, and access information so that clients and servers can share state information to manage use and access.
- Since files are the main information-storage mechanism in most computer systems, *file protection* is needed on multiuser systems. Access to files can be controlled separately for each type of access — read, write, execute, append, delete, list directory, and so on. File protection can be provided by access lists, passwords, or other techniques.



End of Chapter 10.A.

